Open in Colab

Copyright 2023 The MediaPipe Authors.

Licensed under the Apache License, Version 2.0 (the "License");

Show code

Train a custom object detection model with MediaPipe Model Maker

In this colab notebook, you'll learn how to use MediaPipe Model Maker to train a custom object detection model to detect dogs.

The Model Maker library uses *transfer learning* to simplify the process of training a TensorFlow Lite model using a custom dataset to use with the MediaPipe Object Detector task.

Preparation

Install the required packages

Start by installing the required packages, including the Model Maker package from the GitHub repo and the pycocotools library you'll use for evaluation.

```
Start coding or generate with AI.
```

```
!pip install --upgrade pip
!pip install 'keras<3.0.0' mediapipe-model-maker
```

```
Collecting colorama (from sacrebleu->tf-models-official<2.16.
  Downloading colorama-0.4.6-py2.py3-none-any.whl.metadata (1
Requirement already satisfied: lxml in /usr/local/lib/python3
Requirement already satisfied: scikit-learn>=0.21.3 in /usr/l
```

Resources

You are subscribed to Colab Pro. Learn more

Available: 143.81 compute units

Usage rate: approximately 0 per hour

You have 0 active sessions.

Manage session

Not connected to runtime.

Change runtime t

```
Requirement already satisfied: docstring-parser<1.0,>=0.15 in
Requirement already satisfied: googleapis-common-protos<2,>=1
Requirement already satisfied: pycparser in /usr/local/lib/py
Requirement already satisfied: proto-plus<2.0.0dev,>=1.22.3 i
Requirement already satisfied: cachetools<6.0,>=2.0.0 in /usr
Requirement already satisfied: requests-oauthlib>=0.7.0 in /u
Requirement already satisfied: joblib>=1.2.0 in /usr/local/li
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/l
Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/loca
Requirement already satisfied: webencodings in /usr/local/lib
Requirement already satisfied: text-unidecode>=1.3 in /usr/lo
Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/
Downloading keras-2.15.0-py3-none-any.whl (1.7 MB)
     ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 1.7/1.7 MB 40.5 MB
Downloading mediapipe_model_maker-0.2.1.4-py3-none-any.whl (1
Downloading mediapipe-0.10.20-cp310-cp310-manylinux_2_28_x86_
     ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 35.6/35.6 MB 107.
Downloading tensorflow-2.15.1-cp310-cp310-manylinux_2_17_x86_
     ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 475.2/475.2 MB 49
Downloading tensorflow_model_optimization-0.7.5-py2.py3-none-
Downloading tf_models_official-2.15.0-py2.py3-none-any.whl (2
     ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 2.7/2.7 MB 61.9 MB
Downloading tensorflow_text-2.15.0-cp310-cp310-manylinux_2_17_
     ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 5.2/5.2 MB 73.5 MB
Downloading tensorflow_addons-0.23.0-cp310-cp310-manylinux_2_
     ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 611.8/611.8 kB 25
Downloading ml_dtypes-0.3.2-cp310-cp310-manylinux_2_17_x86_64
     ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 2.2/2.2 MB 65.1 MB
Downloading sounddevice-0.5.1-py3-none-any.whl (32 kB)
Downloading tensorboard-2.15.2-py3-none-any.whl (5.5 MB)
     ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 5.5/5.5 MB 90.3 MB
Downloading tensorflow_estimator-2.15.0-py2.py3-none-any.whl
Downloading tf_keras-2.15.1-py3-none-any.whl (1.7 MB)
     ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 1.7/1.7 MB 61.2 MB
Downloading typeguard-2.13.3-py3-none-any.whl (17 kB)
Downloading wrapt-1.14.1-cp310-cp310-manylinux_2_5_x86_64.many
Downloading sacrebleu-2.4.3-py3-none-any.whl (103 kB)
Downloading colorama-0.4.6-py2.py3-none-any.whl (25 kB)
Downloading portalocker-3.0.0-py3-none-any.whl (19 kB)
Building wheels for collected packages: seqeval
  Building wheel for seqeval (setup.py) ... done
  Created wheel for seqeval: filename=seqeval-1.2.2-py3-none-
  Stored in directory: /root/.cache/pip/wheels/1a/67/4a/ad408
Successfully built seqeval
Installing collected packages: wrapt, typeguard, tensorflow-m
  Attempting uninstall: wrapt
    Found existing installation: wrapt 1.17.0
    Uninstalling wrapt-1.17.0:
      Successfully uninstalled wrapt-1.17.0
  Attempting uninstall: typeguard
    Found existing installation: typeguard 4.4.1
    Uninstalling typeguard-4.4.1:
      Successfully uninstalled typeguard-4.4.1
```

```
      Attempting uninstall: ml-dtypes
```

Import the required packages.

```
import os
import tensorflow as tf
assert tf.__version__.startswith('2')
from google.colab import files

from mediapipe_model_maker import object_detector
```

```
/usr/local/lib/python3.10/dist-packages/tensorflow_addons/uti

    TensorFlow Addons (TFA) has ended development and introductio
    TFA has entered a minimal maintenance and release mode until
    Please modify downstream libraries to take dependencies from

    For more information see: https://github.com/tensorflow/addon

      warnings.warn(
```

Prepare the dataset

This dataset contains about 130 images of 3 different dogs: a Golden Retriever, a Labrador Retriever mix, and a brindle-colored Formosan Mountain Dog. This is an example image from the dataset.



We start with downloading the dataset.

```
#!wget https://storage.googleapis.com/mediapipe-assets/dogs2.zip
!unzip rocket8.zip
train_dataset_path = "rocket8/train"
validation_dataset_path = "rocket8/validate"
```

```
  inflating: rocket8/train/Annotations/img301.xml
```

```
inflating: __MACOSX/rocket8/train/Annotations/._img301.xml
inflating: rocket8/train/Annotations/img329.xml
inflating: __MACOSX/rocket8/train/Annotations/._img329.xml
inflating: rocket8/train/Annotations/img117.xml
inflating: __MACOSX/rocket8/train/Annotations/._img117.xml
inflating: rocket8/train/Annotations/img103.xml
inflating: __MACOSX/rocket8/train/Annotations/._img103.xml
inflating: rocket8/train/Annotations/img129.xml
inflating: __MACOSX/rocket8/train/Annotations/._img129.xml
inflating: rocket8/train/Annotations/img115.xml
inflating: __MACOSX/rocket8/train/Annotations/._img115.xml
inflating: rocket8/train/Annotations/img101.xml
inflating: __MACOSX/rocket8/train/Annotations/._img101.xml
inflating: rocket8/train/Annotations/img459.xml
inflating: __MACOSX/rocket8/train/Annotations/._img459.xml
inflating: rocket8/train/Annotations/img471.xml
inflating: __MACOSX/rocket8/train/Annotations/._img471.xml
inflating: rocket8/train/Annotations/img303.xml
inflating: __MACOSX/rocket8/train/Annotations/._img303.xml
inflating: rocket8/train/Annotations/img465.xml
inflating: __MACOSX/rocket8/train/Annotations/._img465.xml
inflating: rocket8/train/Annotations/img539.xml
inflating: __MACOSX/rocket8/train/Annotations/._img539.xml
inflating: rocket8/train/Annotations/img505.xml
inflating: __MACOSX/rocket8/train/Annotations/._img505.xml
inflating: rocket8/train/Annotations/img263.xml
inflating: __MACOSX/rocket8/train/Annotations/._img263.xml
inflating: rocket8/train/Annotations/img277.xml
inflating: __MACOSX/rocket8/train/Annotations/._img277.xml
inflating: rocket8/train/Annotations/img511.xml
inflating: __MACOSX/rocket8/train/Annotations/._img511.xml
inflating: rocket8/train/Annotations/img288.xml
inflating: __MACOSX/rocket8/train/Annotations/._img288.xml
inflating: rocket8/train/Annotations/img289.xml
inflating: __MACOSX/rocket8/train/Annotations/._img289.xml
inflating: rocket8/train/Annotations/img276.xml
inflating: __MACOSX/rocket8/train/Annotations/._img276.xml
inflating: rocket8/train/Annotations/img510.xml
inflating: __MACOSX/rocket8/train/Annotations/._img510.xml
inflating: rocket8/train/Annotations/img504.xml
inflating: __MACOSX/rocket8/train/Annotations/._img504.xml
inflating: rocket8/train/Annotations/img262.xml
inflating: __MACOSX/rocket8/train/Annotations/._img262.xml
inflating: rocket8/train/Annotations/img538.xml
inflating: __MACOSX/rocket8/train/Annotations/._img538.xml
inflating: rocket8/train/Annotations/img13.xml
inflating: __MACOSX/rocket8/train/Annotations/._img13.xml
inflating: rocket8/train/Annotations/img302.xml
inflating: __MACOSX/rocket8/train/Annotations/._img302.xml
inflating: rocket8/train/Annotations/img464.xml
inflating: __MACOSX/rocket8/train/Annotations/._img464.xml
inflating: rocket8/train/Annotations/img470.xml
inflating: __MACOSX/rocket8/train/Annotations/._img470.xml
```

```
      inflating: rocket8/train/Annotations/img316.xml
      inflating: __MACOSX/rocket8/train/Annotations/._img316.xml
      inflating: rocket8/train/Annotations/img458.xml
      inflating: __MACOSX/rocket8/train/Annotations/._img458.xml
      inflating: rocket8/train/Annotations/img114.xml
```

Train the object detection model

Step 1: Load the dataset

- Images in `train_data` is used to train the custom object detection model.
- Images in `val_data` is used to check if the model can generalize well to new images that it hasn't seen before.

```
train_data = object_detector.Dataset.from_pascal_voc_folder(
    'rocket8/train',
    cache_dir="/tmp/od_data/train",
)

val_data = object_detector.Dataset.from_pascal_voc_folder(
    'rocket8/validate',
    cache_dir="/tmp/od_data/validatation")
```

Step 2: Train the TensorFlow model with the training data.

- Set `batch_size = 8` here so you will see that it takes 13 steps to go through the 139 images in the training dataset.
- Set `learning_rate = 0.3` here so the model will adjust its weights at a rate of 0.3.
- Set `epochs = 50`, which means it will go through the training dataset 50 times. You can look at the validation accuracy during training and stop when you see validation loss (`val_loss`) stop decreasing to avoid overfitting.

```
hparams = object_detector.HParams(batch_size=8, learning_rate=0.3,
options = object_detector.ObjectDetectorOptions(
    supported_model=object_detector.SupportedModels.MOBILENET_V2,
    hparams=hparams
)
model = object_detector.ObjectDetector.create(
```

```
train_data=train_data,
validation_data=val_data,
options=options)
```

```
/usr/local/lib/python3.10/dist-packages/keras/src/engine/funct
  inputs = self._flatten_to_reference_inputs(inputs)
WARNING:tensorflow:`tf.keras.layers.experimental.SyncBatchNorm
Downloading https://storage.googleapis.com/tf_model_garden/vis
Model: "retina_net_model"
_________________________________________________________________
 Layer (type)                Output Shape              Param #
=================================================================
 mobile_net (MobileNet)      {'2': (None, 64, 64, 24   2257984
                             ),
                              '3': (None, 32, 32, 32
                             ),
                              '4': (None, 16, 16, 96
                             ),
                              '5': (None, 8, 8, 320)
                             , '6': (None, 8, 8, 128
                             0)}

 fpn (FPN)                   {'5': (None, 8, 8, 128)   149056
                             , '4': (None, 16, 16, 1
                             28),
                              '3': (None, 32, 32, 12
                             8),
                              '6': (None, 4, 4, 128)
                             , '7': (None, 2, 2, 128
                             )}

 multilevel_detection_gener  multiple                  0 (unus
 ator (MultilevelDetectionG
 enerator)

 retina_net_head (RetinaNet  ({'3': (None, 32, 32, 8   179189
 Head)                       1),
                              '4': (None, 16, 16, 81
                             ),
                              '5': (None, 8, 8, 81),
                              '6': (None, 4, 4, 81),
                              '7': (None, 2, 2, 81)}
                             , {'3': (None, 32, 32,
                             36),
                              '4': (None, 16, 16, 36
                             ),
                              '5': (None, 8, 8, 36),
                              '6': (None, 4, 4, 36),
                              '7': (None, 2, 2, 36)}
                             , {})

=================================================================
Total params: 2586229 (9.87 MB)
```

```
Trainable params: 2540597 (9.69 MB)
Non-trainable params: 45632 (178.25 KB)
_________________________________________________________________
Epoch 1/50
/usr/local/lib/python3.10/dist-packages/keras/src/backend.py:
  warnings.warn(
WARNING:tensorflow:Gradients do not exist for variables ['con
WARNING:tensorflow:Gradients do not exist for variables ['con
WARNING:tensorflow:Gradients do not exist for variables ['con
WARNING:tensorflow:Gradients do not exist for variables ['con
```

Step 3. Evaluate the model with the validation data.

After training the object detection model using the images in the training dataset, use the 27 images in the validation dataset to evaluate how the model performs against new data it has never seen before.

The evaluation metrics are same as COCO.

```
loss, coco_metrics = model.evaluate(val_data, batch_size=4)
print(f"Validation loss: {loss}")
print(f"Validation coco metrics: {coco_metrics}")
```

```
17/17 [==============================] - 3s 91ms/step - total_
creating index...
index created!
creating index...
index created!
Running per image evaluation...
Evaluate annotation type *bbox*
DONE (t=0.31s).
Accumulating evaluation results...
DONE (t=0.12s).
 Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxI
 Average Precision  (AP) @[ IoU=0.50      | area=   all | maxI
 Average Precision  (AP) @[ IoU=0.75      | area=   all | maxI
 Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxI
 Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxI
 Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxI
 Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxI
 Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxI
 Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxI
 Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxI
 Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxI
 Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxI
Validation loss: [0.7392598986625671, 0.4557432532310486, 0.0(
Validation coco metrics: {'AP': 0.39894685, 'AP50': 0.8752042
```

Step 4: Export as a TensorFlow Lite model.

After creating the model, convert and export it to a Tensorflow Lite model format, and then download it for later use in an on-device application with MediaPipe Tasks.

```
model.export_model('rocket8.tflite')
!ls exported_model
files.download('exported_model/rocket8.tflite')
```

```
Exporting a floating point model
/usr/local/lib/python3.10/dist-packages/keras/src/engine/funct
  inputs = self._flatten_to_reference_inputs(inputs)
/usr/local/lib/python3.10/dist-packages/keras/src/engine/funct
  inputs = self._flatten_to_reference_inputs(inputs)
WARNING:tensorflow:Skipping full serialization of Keras layer
WARNING:tensorflow:Skipping full serialization of Keras layer
WARNING:tensorflow:Skipping full serialization of Keras layer
WARNING:tensorflow:Skipping full serialization of Keras layer
WARNING:tensorflow:Skipping full serialization of Keras layer
WARNING:tensorflow:Skipping full serialization of Keras layer
WARNING:tensorflow:Skipping full serialization of Keras layer
WARNING:tensorflow:Skipping full serialization of Keras layer
WARNING:tensorflow:Skipping full serialization of Keras layer
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
```

```
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
```

```
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
```

```
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
```

```
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
```

```
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
```

```
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
```

```
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
```

```
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
```

```
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
```

```
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
```

```
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
```

```
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
```

```
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
WARNING:absl:Importing a function (__inference_internal_grad_1
```

```
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
```

```
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
```

```
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
```

```
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
WARNING:absl:Importing a function (__inference_internal_grad_f
dogs.tflite  float_ckpt.data-00000-of-00001  float_ckpt.index
```